

SQL-Focused

You are building a database for ACME website to show users' skills. Each user can have multiple skills (Python, Java, SQL, NodeJS, etc...). You need to design the database schema and write a SQL query to find users that have at least 5 of the top 10 skills.

```
SELECT u.user_id, u.name
FROM User_data u
INNER JOIN UserSkills us ON u.user_id = us.user_id
INNER JOIN (
    SELECT skill_id
    FROM Skills
    ORDER BY skill_name DESC
    LIMIT 10
) top_skills ON us.skill_id = top_skills.skill_id
GROUP BY u.user_id, u.name
HAVING COUNT(*) >= 5;
```

Questions:

- Can you explain the performance of your query?

- The performance of the query depends on the size of the UserSkills and Skills tables, as well as the available indexes.
- It retrieves the top 10 skills based on skill_name and then joins the User and UserSkills tables to find users with at least 5 of those skills
- The performance can be impacted if the tables have a large number of records or if the necessary indexes are not in place.

- How can you improve the query performance?

- Index the skill_name column in the Skills table for efficient sorting.
- Regularly analyze and update statistics to help the query optimizer make informed decisions.
- Consider implementing query caching or result caching if the query is frequently executed with the same parameters.

- Can you suggest different indexes or any other ways to optimize the query execution time?

- Index the skill_name and name column in tables.
- Apart from indexing and query optimization, consider denormalizing the data to avoid joins and improve query performance.
- Store a denormalized list of skills directly in the User table, allowing for faster retrieval without the need for joins.

An example task that we can give to candidates to assess their ability to write queries using window functions
Task: You have a table that contains information about user work experience.

The table has the following columns:

- id (integer)
- user_id (integer)
- title (text)
- description (text)
- date_added (date)
- date_edited (date)

You need to write a SQL query to get only 2nd work experience, for example user King Arthur has experience:

id	user_id	title	description
1000	1	Farmer	Collected potatoes
1001	1	Knight	Killed enemies
1002	1	King	Ruled the kingdom

The result of the query should be a record with id == 1001

```
SELECT *
FROM (
  SELECT *, ROW_NUMBER() OVER (PARTITION BY user_id ORDER BY date_added) AS rn
  FROM WorkExperience
) AS subquery
WHERE rn = 2 AND user_id = <user_id>;
```

Questions:

- Can you explain the performance of your query?

- The performance depends on the size of the WorkExperience table and the availability of indexes.
- The query uses the ROW_NUMBER() window function to assign a row number to each work experience based on the date_added column. It then filters the result to fetch the row with rn = 2.
- The performance can be impacted if the table has a large number of records or if the necessary indexes are not in place.

- How can you improve the query performance?

- Ensure proper indexing on the user_id and date_added columns in the WorkExperience table to optimize sorting and partitioning.
- Regularly analyze and update statistics to help the query optimizer make informed decisions.
- Consider implementing query caching or result caching if the query is frequently executed with the same parameters.

- Can you suggest any other approach to solve this problem?

```
SELECT *
FROM WorkExperience
WHERE user_id = <user_id>
ORDER BY date_added
LIMIT 1 OFFSET 1;
```

PySpark

Environment: Databricks or EMR cluster Requirements: PySpark, Delta Lake, AWS S3

QUESTIONS:

Analyzing Customer Reviews (40 minutes) You have been provided with a large dataset of customer reviews in Delta Lake format on AWS S3: `s3://your-bucket/customer_reviews/`

The dataset contains a mixture of structured and unstructured data. Each line in the dataset represents a customer review, and has the following format in a set of delta files: `product_id`, `user_id`, `user_name`, `rating`, `review_date`, `review_title`, `review_text`

Your task is to write a PySpark script that reads this dataset and performs the following analysis: However, due to the presence of messy data, some lines may contain extra or missing fields,

improperly formatted dates, or special characters.

1. Create a table (format choice is yours) with the following columns: `uuid` (String): Unique identifier for each record, generated based on `product_id`, `user_id`, and the Unix timestamp of the review `product_id` (String): Unique identifier for each product `user_id` (String): Unique identifier for each user `user_name` (String): Name of the user `rating` (Float): Rating given by the user, ranging from 1.0 to 5.0 `review_date` (Date): Date of the review submission `review_title` (String): Title of the review `review_text` (String): Text of the review

2. Read the dataset and validate the data by: Validating each column using regex patterns, as applicable Removing any special characters from the `user_name`, `review_title`, and `review_text` columns Filtering out rows with an incorrect number of fields. If missing more than 40% of data. Converting improperly formatted dates in the `review_date` column to a standard date format (e.g., "yyyy-MM-dd") Alerting if a row does not match the regex patterns or contains an improperly formatted date (only if you dont clean the data col) Remove columns if column is missing more than 90% of data.

3. Calculate the average rating for each product.

4. Find the top 10 products with the highest average ratings.

5. Identify the top three users with the most reviews submitted. Please use appropriate PySpark APIs and follow best practices for code readability and optimization. Ensure that your solution is optimized for processing large datasets (100 TBs) by leveraging big data best practices.

```

from pyspark.sql import SparkSession
from pyspark.sql.functions import col, udf, from_unixtime, current_date
from pyspark.sql.types import StringType, FloatType, DateType
import re

# Create SparkSession
spark = SparkSession.builder.appName("CustomerReviewsAnalysis").getOrCreate()

# Step 1: Create a table with the specified columns
spark.sql("""
    CREATE TABLE IF NOT EXISTS customer_reviews (
        uuid STRING,
        product_id STRING,
        user_id STRING,
        user_name STRING,
        rating FLOAT,
        review_date DATE,
        review_title STRING,
        review_text STRING
    )
    USING DELTA
    PARTITIONED BY (product_id)
    LOCATION 's3://your-bucket/customer_reviews/'
""")

# Step 2: Read the dataset and validate the data
df = spark.read.format("delta").load("s3://your-bucket/customer_reviews/")

# Define regex patterns for validation
uuid_pattern = re.compile(r"^[w{8}-w{4}-w{4}-w{4}-w{12}$")
date_pattern = re.compile(r"^[d{4}-d{2}-d{2}$")

# UDF for generating UUID based on product_id, user_id, and review_date
generate_uuid = udf(lambda product_id, user_id, review_date: str(hash(product_id +
user_id + str(review_date))))

# UDF for cleaning special characters from strings
clean_string = udf(lambda s: re.sub(r"^[w\s]", "", s) if s is not None else None)

# UDF for converting improperly formatted dates to standard date format
convert_date = udf(lambda d: None if d is None else d if date_pattern.match(d) else
None, DateType())

# Apply data validation and cleaning
df = df.withColumn("uuid", generate_uuid(col("product_id"), col("user_id"),
col("review_date"))) \
    .withColumn("user_name", clean_string(col("user_name"))) \
    .withColumn("review_title", clean_string(col("review_title"))) \
    .withColumn("review_text", clean_string(col("review_text"))) \
    .withColumn("review_date", convert_date(col("review_date"))) \
    .filter(col("uuid").rlike(uuid_pattern)) \
    .filter(col("rating").between(1.0, 5.0)) \
    .filter(col("product_id").isNotNull() & col("user_id").isNotNull() &
col("user_name").isNotNull() &
        col("rating").isNotNull() & col("review_date").isNotNull() &
col("review_title").isNotNull() &
        col("review_text").isNotNull()) \
    .filter(col("user_name") != "") \
    .filter(col("review_title") != "") \
    .filter(col("review_text") != "") \
    .filter(col("review_date").isNotNull()) \
    .filter(col("review_date") <= current_date())

```

```
# Step 3: Calculate average rating for each product
average_rating = df.groupBy("product_id").agg({"rating": "avg"})

# Step 4: Find top 10 products with highest average ratings
top_products = average_rating.orderBy(col("avg(rating)").desc()).limit(10)

# Step 5: Identify top three users with most reviews submitted
top_users = df.groupBy("user_id").count().orderBy(col("count").desc()).limit(3)

# Display the results
top_products.show()
top_users.show()

# Stop the SparkSession
spark
```

Data Migration



You need to perform the migration of an Oracle database to Snowflake. Your goal is to efficiently transfer the data from Oracle to Snowflake and ensure that the data is stored correctly in the new environment.

Considerations

- Import the database from the link below into your Oracle database
- Analysis of the Oracle database
- Schema and table mapping
- Extraction of data from Oracle
- Creation of tables in Snowflake
- Data loading into Snowflake using Informatica
- Migration of database logic and application code
- Testing and validation
- Documentation

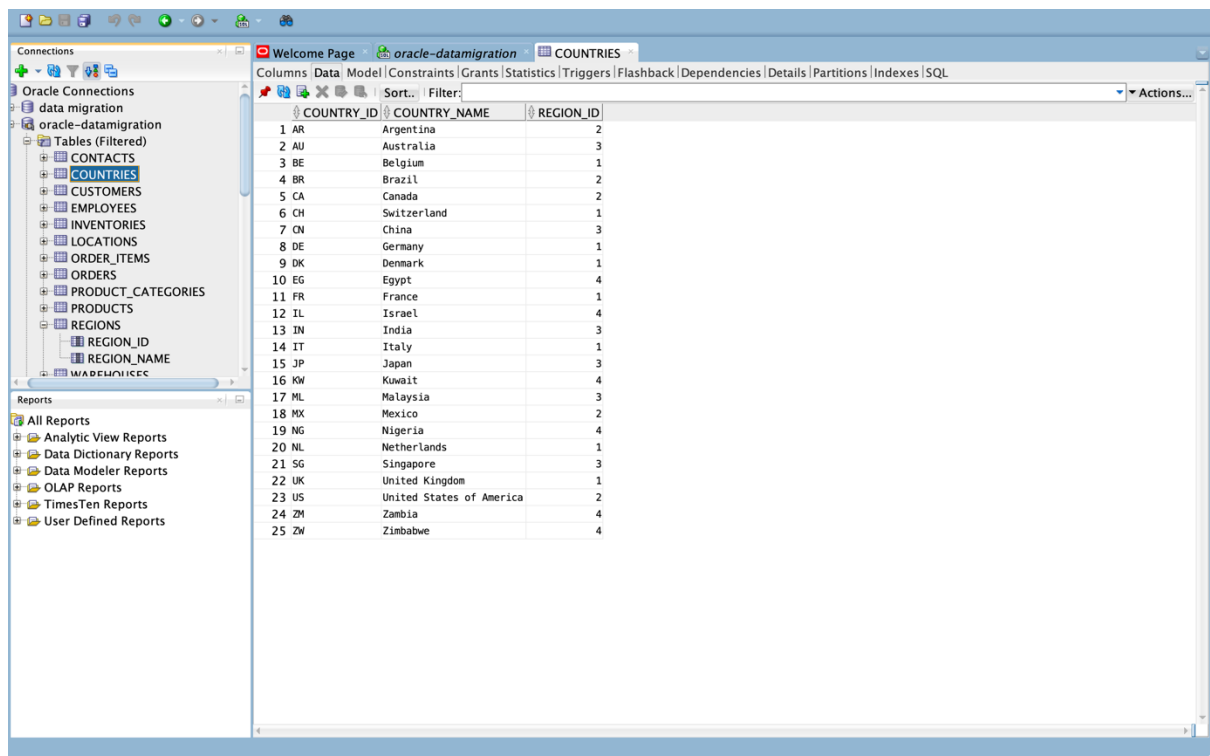
The provided code should not be a script, it should be able to reuse cross-projects or applications. Be sure that the code you sent is correctly tested and well documented. The database to be migrated <https://www.oracletutorial.com/wp-content/uploads/2019/01/oracle-sample-database.zip>

- Import the database from the link below into your Oracle database
- Obtain the Oracle setup and import scripts from the provided link.

The screenshot shows the Oracle SQL Developer interface. The 'Connections' pane on the left shows the 'oracle-datamigration' connection selected. The 'Tables' pane on the right displays a list of tables in the 'oracle-datamigration' schema. The table list includes columns for TABLE_NAME, NUM_ROWS, STATUS, COLUMNS, COMMENTS, INDEXED_COLUMNS, AVG_ROW_LEN, TABLESPACE_NAME, and ESTIMATED_SIZE. The tables listed are: 1 CONTACTS, 2 COUNTRIES, 3 CUSTOMERS, 4 EMPLOYEES, 5 INVENTORIES, 6 LOCATIONS, 7 ORDERS, 8 ORDER_ITEMS, 9 PRODUCTS, 10 PRODUCT_CATEGORIES, 11 REGIONS, and 12 WAREHOUSES.

TABLE_NAME	NUM_ROWS	STATUS	COLUMNS	COMMENTS	INDEXED_COLUMNS	AVG_ROW_LEN	TABLESPACE_NAME	ESTIMATED_SIZE	LAS
1 CONTACTS	(null)	VALID	6 (null)		1	(null)	USERS	(null)	08-06
2 COUNTRIES	(null)	VALID	3 (null)		1	(null)	USERS	(null)	08-06
3 CUSTOMERS	(null)	VALID	5 (null)		1	(null)	USERS	(null)	08-06
4 EMPLOYEES	(null)	VALID	8 (null)		1	(null)	USERS	(null)	08-06
5 INVENTORIES	(null)	VALID	3 (null)		1	(null)	USERS	(null)	08-06
6 LOCATIONS	(null)	VALID	6 (null)		1	(null)	USERS	(null)	08-06
7 ORDERS	(null)	VALID	5 (null)		1	(null)	USERS	(null)	08-06
8 ORDER_ITEMS	(null)	VALID	5 (null)		1	(null)	USERS	(null)	08-06
9 PRODUCTS	(null)	VALID	6 (null)		1	(null)	USERS	(null)	08-06
10 PRODUCT_CATEGORIES	(null)	VALID	2 (null)		1	(null)	USERS	(null)	08-06
11 REGIONS	(null)	VALID	2 (null)		1	(null)	USERS	(null)	08-06
12 WAREHOUSES	(null)	VALID	3 (null)		1	(null)	USERS	(null)	08-06

- Analysis of the Oracle database



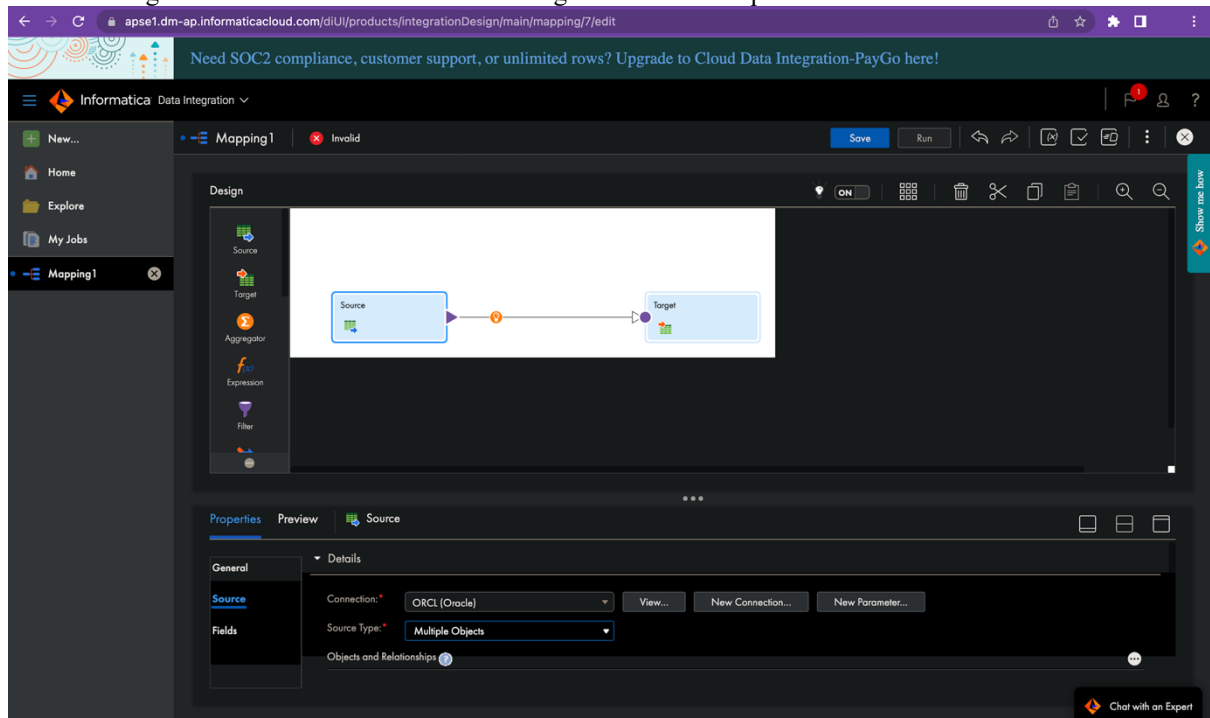
The screenshot shows the Oracle SQL Developer interface. On the left, the 'Connections' pane shows 'oracle-datamigration' selected. Under 'Tables (Filtered)', the 'COUNTRIES' table is highlighted. The main window displays the 'COUNTRIES' table data with columns: COUNTRY_ID, COUNTRY_NAME, and REGION_ID. The data is sorted by COUNTRY_ID.

COUNTRY_ID	COUNTRY_NAME	REGION_ID
1 AR	Argentina	2
2 AU	Australia	3
3 BE	Belgium	1
4 BR	Brazil	2
5 CA	Canada	2
6 CH	Switzerland	1
7 CN	China	3
8 DE	Germany	1
9 DK	Denmark	1
10 EG	Egypt	4
11 FR	France	1
12 IL	Israel	4
13 IN	India	3
14 IT	Italy	1
15 JP	Japan	3
16 KW	Kuwait	4
17 ML	Malaysia	3
18 MX	Mexico	2
19 NG	Nigeria	4
20 NL	Netherlands	1
21 SG	Singapore	3
22 UK	United Kingdom	1
23 US	United States of America	2
24 ZH	Zambia	4
25 ZW	Zimbabwe	4

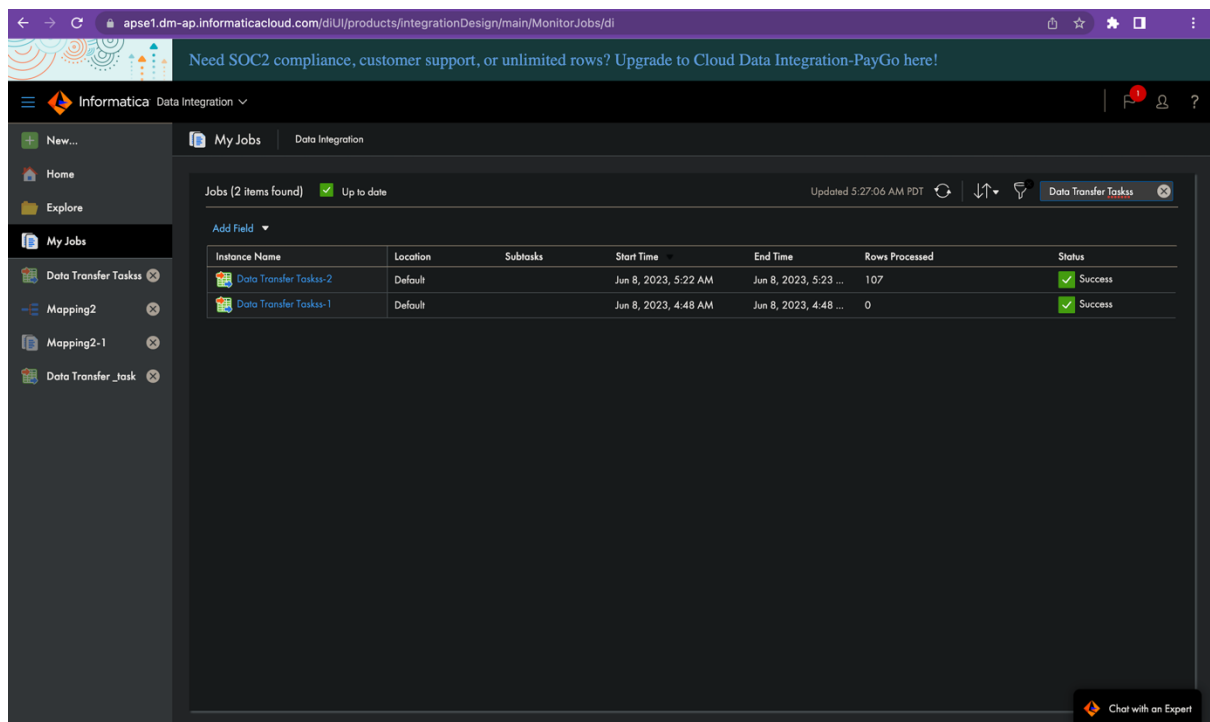
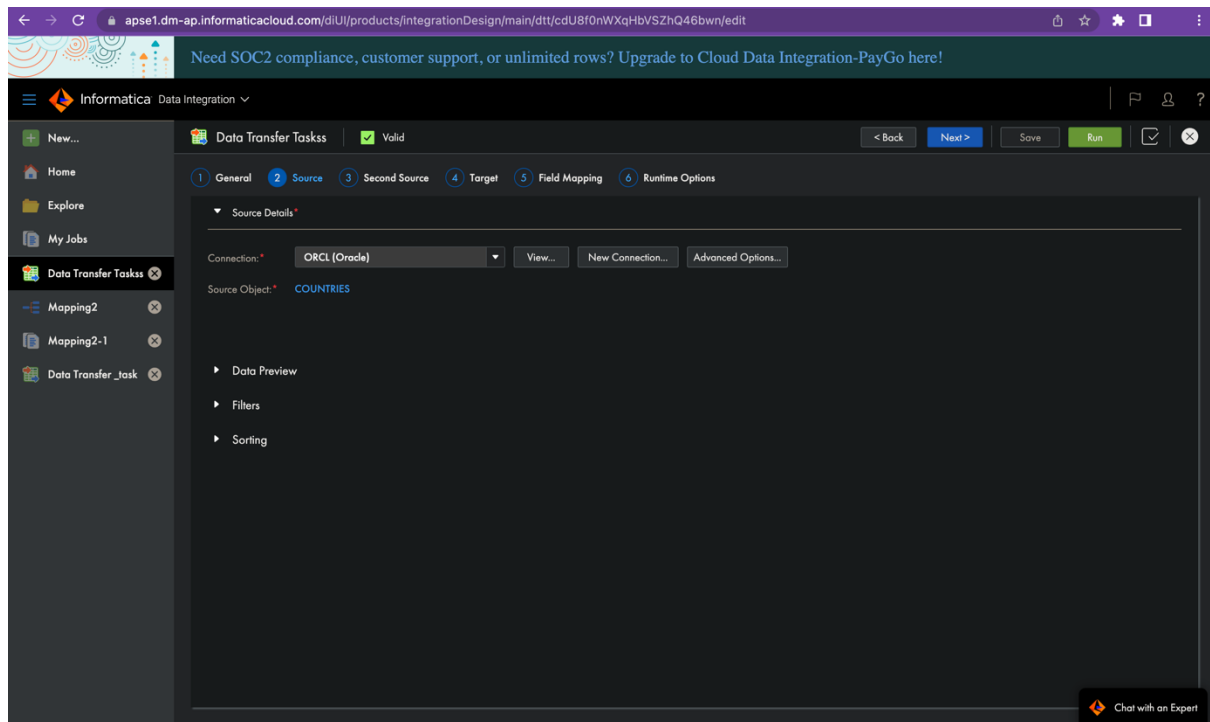
- Schema and table mapping

In Informatica, the procedure involves configuring both the oracle as a source and snowflake as a destination database and establishing a mapping that links the two.

We can change the source and destination according to our further requirements.



- Extraction of data from Oracle



- Creation of tables in Snowflake

To create tables in Snowflake, I wrote a snowsql script.

The screenshot shows the Snowflake web interface with a SQL script in the editor. The script is for creating three tables: regions, countries, and locations. The regions table has columns region_id (primary key) and region_name. The countries table has columns country_id (primary key), country_name, and region_id (foreign key to regions). The locations table has columns location_id (primary key) and region_id (foreign key to regions). The script is executed, and the results show 0 rows.

```
7 CREATE TABLE regions
8 (
9   region_id NUMBER DEFAULT regions_seq.NEXTVAL PRIMARY KEY,
10  region_name VARCHAR(50) NOT NULL
11 );
12
13 -- countries table
14 CREATE TABLE countries
15 (
16  country_id CHAR(2) PRIMARY KEY,
17  country_name VARCHAR(40) NOT NULL,
18  region_id NUMBER, -- fk
19  CONSTRAINT fk_countries_regions FOREIGN KEY (region_id)
20    REFERENCES regions(region_id)
21    ON DELETE CASCADE
22 );
23
24 -- locations table
25 CREATE TABLE locations
26 (
27  location_id NUMBER AUTOINCREMENT() PRIMARY KEY,
```

Query Details

Query duration	26ms
Rows	0
Query ID	01accda3-3200-c0fc-0...

The screenshot shows the Snowflake web interface with a SQL script in the editor. The script is for creating three tables: warehouses, employees, and product_categories. The warehouses table has columns warehouse_id (primary key) and warehouse_name. The employees table has columns employee_id (primary key), first_name, last_name, email, phone, hire_date, and manager_id (foreign key to employees). The product_categories table has columns product_id (primary key) and product_name. The script is executed, and the results show 0 rows.

```
24 -- warehouses table
25 CREATE TABLE warehouses
26 (
27  warehouse_id NUMBER AUTOINCREMENT() PRIMARY KEY,
28  warehouse_name VARCHAR(255),
29  location_id NUMBER, -- fk
30  CONSTRAINT fk_warehouses_locations FOREIGN KEY (location_id)
31    REFERENCES locations(location_id)
32    ON DELETE CASCADE
33 );
34
35 -- employees table
36 CREATE TABLE employees
37 (
38  employee_id NUMBER AUTOINCREMENT() PRIMARY KEY,
39  first_name VARCHAR(255) NOT NULL,
40  last_name VARCHAR(255) NOT NULL,
41  email VARCHAR(255) NOT NULL,
42  phone VARCHAR(50) NOT NULL,
43  hire_date DATE NOT NULL,
44  manager_id NUMBER, -- fk
45  job_title VARCHAR(255) NOT NULL,
46  CONSTRAINT fk_employees_manager FOREIGN KEY (manager_id)
47    REFERENCES employees(employee_id)
48    ON DELETE CASCADE
49 );
50
51 -- product_categories table
52 CREATE TABLE product_categories
53 (
54  product_id NUMBER AUTOINCREMENT() PRIMARY KEY,
```

Query Details

Query duration	26ms
Rows	0
Query ID	01accda3-3200-c0fc-0...

test_data_migration - Snowflake

app.snowflake.com/ap-southeast-1/he94254/whjOzbQfzia/query

Worksheets target_node 2023-06-07 11:11am data migration test_data_migration PREVIEW

Databases Worksheets

Search

Benchmarking Tutorials

test_data_migration

data migration

2023-06-07 11:11am

target_node

DB_DATA_MIGRATION.PUBLIC Settings Latest Version

```
-- product_categories table
65 CREATE TABLE product_categories
66 (
67   category_id NUMBER AUTOINCREMENT() PRIMARY KEY,
68   category_name VARCHAR(255) NOT NULL
69 );
70
71
72 -- products table
73 CREATE TABLE products
74 (
75   product_id NUMBER AUTOINCREMENT() PRIMARY KEY,
76   product_name VARCHAR(255) NOT NULL,
77   description VARCHAR(2000),
78   standard_cost NUMBER(9, 2),
79   list_price NUMBER(9, 2),
80   category_id NUMBER NOT NULL,
81   CONSTRAINT fk_products_categories FOREIGN KEY (category_id)
82     REFERENCES product_categories(category_id)
83     ON DELETE CASCADE
84 );
85
86 -- customers table
87 CREATE TABLE customers
88 (
89   customer_id NUMBER AUTOINCREMENT() PRIMARY KEY,
90   name VARCHAR(255) NOT NULL,
91   address VARCHAR(255),
92   website VARCHAR(255),
93   credit_limit NUMBER(8, 2)
94 );
95
96 -- contacts table
97 CREATE TABLE contacts
98 (
99   contact_id NUMBER AUTOINCREMENT() PRIMARY KEY,
100   first_name VARCHAR(255) NOT NULL,
101   last_name VARCHAR(255) NOT NULL,
102   email VARCHAR(255) NOT NULL,
103   phone VARCHAR(20),
104   customer_id NUMBER,
105   CONSTRAINT fk_contacts_customers FOREIGN KEY (customer_id)
106     REFERENCES customers(customer_id)
107     ON DELETE CASCADE
108 );
```

test_data_migration - Snowflake

app.snowflake.com/ap-southeast-1/he94254/whjOzbQfzia/query

Worksheets target_node 2023-06-07 11:11am data migration test_data_migration PREVIEW

Databases Worksheets

Search

Benchmarking Tutorials

test_data_migration

data migration

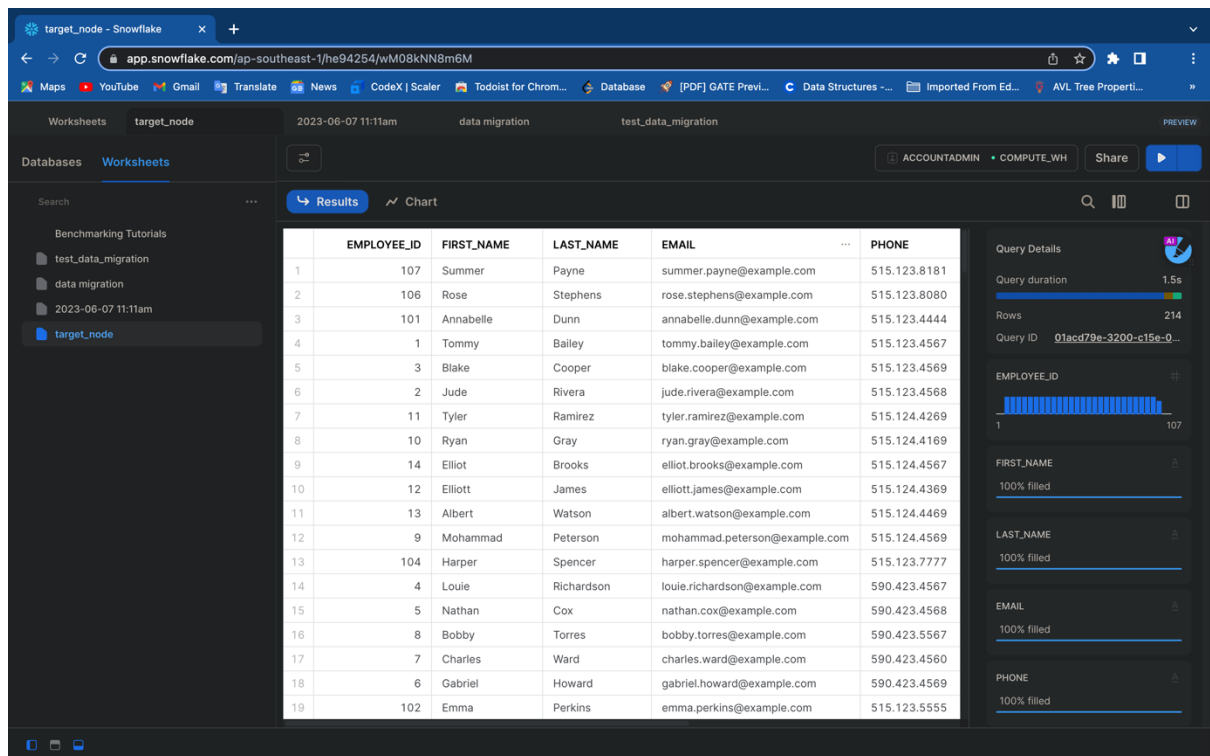
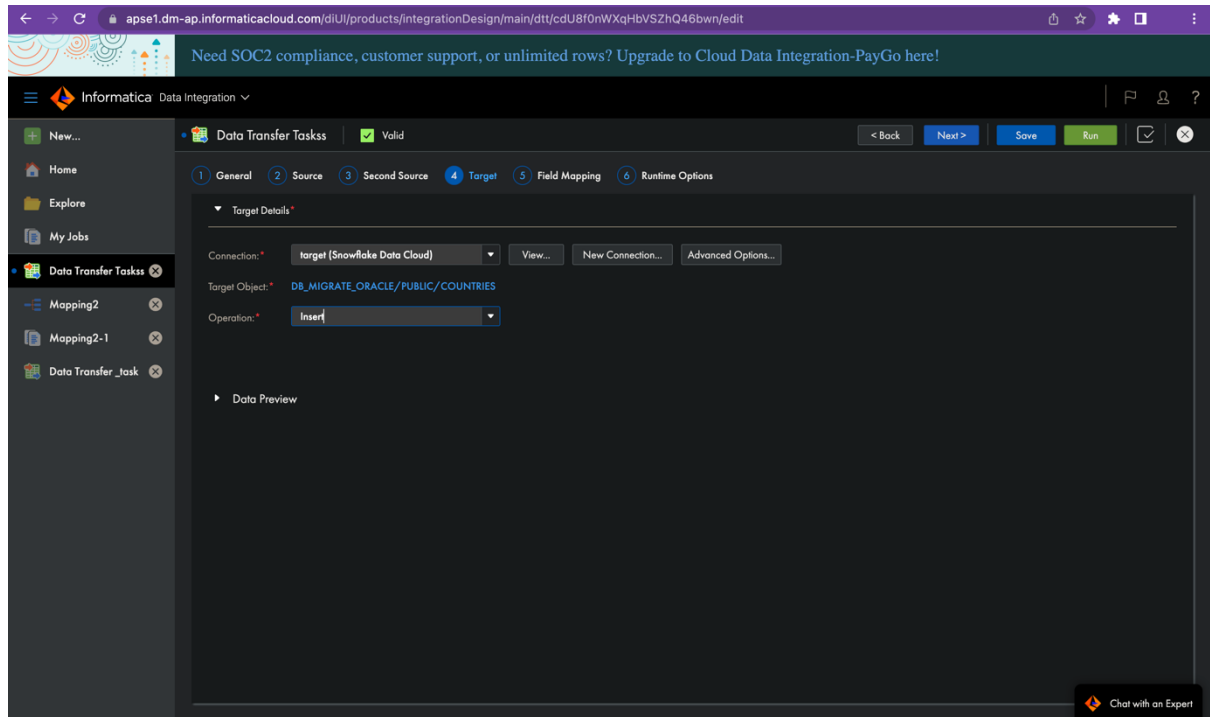
2023-06-07 11:11am

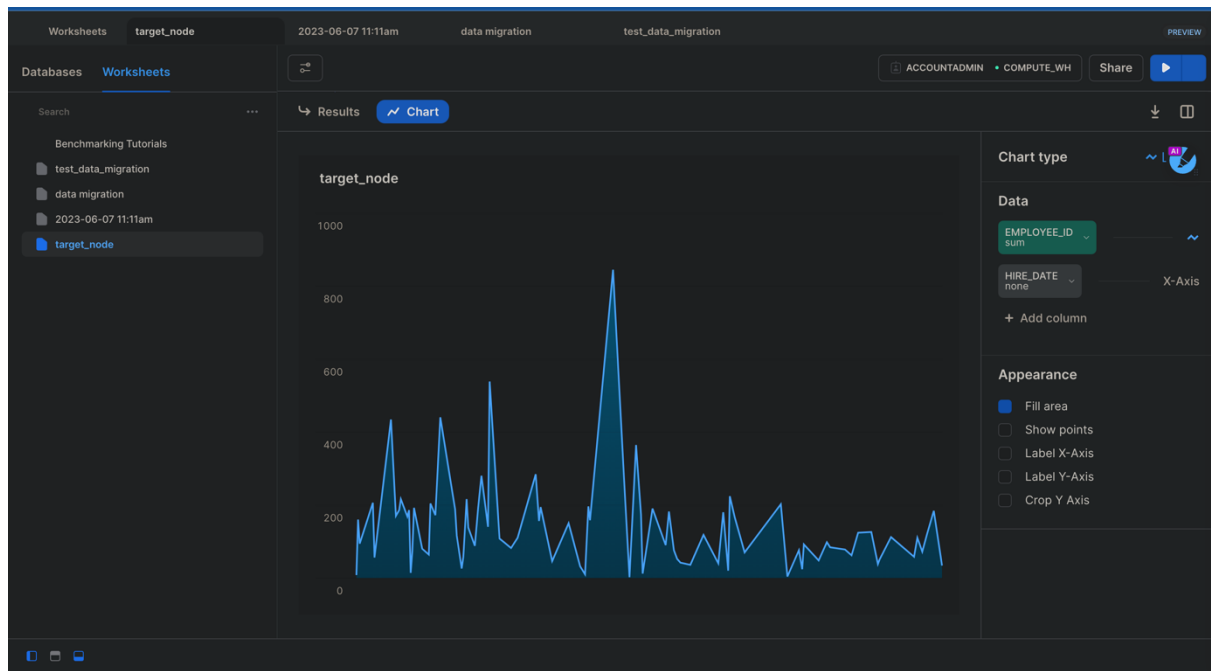
target_node

DB_DATA_MIGRATION.PUBLIC Settings Latest Version

```
113 order_id NUMBER AUTOINCREMENT() PRIMARY KEY,
114 customer_id NUMBER NOT NULL, -- fk
115 status VARCHAR(20) NOT NULL,
116 salesman_id NUMBER, -- fk
117 order_date DATE NOT NULL,
118 CONSTRAINT fk_orders_customers FOREIGN KEY (customer_id)
119   REFERENCES customers(customer_id)
120   ON DELETE CASCADE,
121 CONSTRAINT fk_orders_employees FOREIGN KEY (salesman_id)
122   REFERENCES employees(employee_id)
123   ON DELETE SET NULL
124 );
125
126 -- order_items table
127 CREATE TABLE order_items
128 (
129   order_id NUMBER(12, 0), -- fk
130   item_id NUMBER(12, 0),
131   product_id NUMBER(12, 0) NOT NULL, -- fk
132   quantity NUMBER(8, 2) NOT NULL,
133   unit_price NUMBER(8, 2) NOT NULL,
134   CONSTRAINT pk_order_items PRIMARY KEY (order_id, item_id),
135   CONSTRAINT fk_order_items_products FOREIGN KEY (product_id)
136     REFERENCES products (product_id)
137     ON DELETE CASCADE,
138   CONSTRAINT fk_order_items_orders FOREIGN KEY (order_id)
139     REFERENCES orders (order_id)
140     ON DELETE CASCADE
141 );
142
143 -- inventories table
144 CREATE TABLE inventories
145 (
146   product_id NUMBER(12, 0), -- fk
147   warehouse_id NUMBER(12, 0), -- fk
148   quantity NUMBER(8, 0) NOT NULL,
149   CONSTRAINT pk_inventories PRIMARY KEY (product_id, warehouse_id),
150   CONSTRAINT fk_inventories_products FOREIGN KEY (product_id)
151     REFERENCES products (product_id)
152     ON DELETE CASCADE,
153   CONSTRAINT fk_inventories_warehouses FOREIGN KEY (warehouse_id)
154     REFERENCES warehouses (warehouse_id)
155     ON DELETE CASCADE
156 );
```

- Data loading into Snowflake using Informatica





- Migration of database logic and application code

Need SOC2 compliance, customer support, or unlimited rows? Upgrade to Cloud Data Integration-PayGo here!

Informatica Data Integration

My Jobs

Jobs (10) Up to date Updated 8:53:28 PM PDT

Instance Name	Location	Subtasks	Start Time	End Time	Rows Processed	Status
Data Transfer Taskss-3	Default		Jun 8, 2023, 5:28 AM	Jun 8, 2023, 5:28 ...	107	Success
Data Transfer Taskss-2	Default		Jun 8, 2023, 5:22 AM	Jun 8, 2023, 5:23 ...	107	Success
Data Transfer Taskss-1	Default		Jun 8, 2023, 5:06 AM	Jun 8, 2023, 5:07 ...	0	Success
Mapping2-1	Default		Jun 8, 2023, 5:00 AM	Jun 8, 2023, 5:00 ...	0	Success
Data Transfer Taskss-1	Default		Jun 8, 2023, 4:48 AM	Jun 8, 2023, 4:48 ...	0	Success
Data Transfer Taskss-4-3	Default		Jun 8, 2023, 4:41 AM	Jun 8, 2023, 4:42 ...	0	Success
Data Transfer Taskss-4-2	Default		Jun 8, 2023, 4:38 AM	Jun 8, 2023, 4:39 ...	0	Success
Data Transfer Taskss-4-1	Default		Jun 8, 2023, 4:37 AM	Jun 8, 2023, 4:37 ...	0	Success
Data Transfer Taskss-2-1	Default		Jun 8, 2023, 2:26 AM	Jun 8, 2023, 2:27 ...	1	Success
Data Transfer ETL-1	Default		Jun 7, 2023, 11:17 PM	Jun 7, 2023, 11:17 ...	1	Success

Expert is Offline

Need SOC2 compliance, customer support, or unlimited rows? Upgrade to Cloud Data Integration-PayGo here!

Informatica Data Integration

Data Transfer Taskss-2

Job Properties		Results	
Task Name:	Data Transfer Taskss	Status:	Success
Instance ID:	2	Success Rows:	107
Task Type:	Data Transfer Task	Errors:	0
Started By:	aari.b@consultadd.com through UI	Session Log:	Download Session Log
Start Time:	Jun 8, 2023 5:22:33 AM		
End Time:	Jun 8, 2023 5:23:15 AM		
Duration:	42 seconds		
Runtime Environment:	Informatica Cloud Hosted Agent		
Context Name:	Data Transfer Taskss		

Individual Source/Target Results

Name	Success Rows	Errors	Error Message
Source	107	0	
Target (EMPLOYEES)	107	0	

- Testing and validation

The screenshot shows the Snowflake web interface. The top navigation bar includes links to Maps, YouTube, Gmail, Translate, News, CodeX | Scaler, and Todoist for Chrom... The main header shows the current workspace 'target_node' and the date '2023-06-07 11:11am'. The left sidebar lists databases and schemas, including 'DB_DATA_MIGRATION', 'DB_MIGRATE_ORACLE', 'INFORMATION_SCHEMA', 'PUBLIC', 'SNOWFLAKE', 'SNOWFLAKE_SAMPLE_DATA', and 'TEST_SF'. The main area displays the query 'select * from employees;' and its results in a table with columns: .OYEE_ID, FIRST_NAME, LAST_NAME, EMAIL, and PHONE. The results table contains 12 rows of employee data. The right-hand panel shows 'Query Details' with a duration of 610ms, 214 rows, and a query ID. It also includes a chart for 'EMPLOYEE_ID' and a progress bar for 'FIRST_NAME'.

.OYEE_ID	FIRST_NAME	LAST_NAME	EMAIL	PHONE	
1	107	Summer	Payne	summer.payne@example.com	515.123.8181
2	106	Rose	Stephens	rose.stephens@example.com	515.123.8080
3	101	Annabelle	Dunn	annabelle.dunn@example.com	515.123.4444
4	1	Tommy	Bailey	tommy.bailey@example.com	515.123.4567
5	3	Blake	Cooper	blake.cooper@example.com	515.123.4569
6	2	Jude	Rivera	jude.rivera@example.com	515.123.4568
7	11	Tyler	Ramirez	tyler.ramirez@example.com	515.124.4269
8	10	Ryan	Gray	ryan.gray@example.com	515.124.4169
9	14	Elliot	Brooks	elliott.brooks@example.com	515.124.4567
10	12	Elliott	James	elliott.james@example.com	515.124.4369
11	13	Albert	Watson	albert.watson@example.com	515.124.4469
12	9	Mohammad	Peterson	mohammad.peterson@example.com	515.124.4569

Logs:

```

Agent Group Id: 02500000000002
Agent Group Name: Informatica Cloud Hosted Agent
Agent Id: 00800000000003Q
Task Name: Data Transfer Taskss
Task Type: DTT
Task Id: 0123UT49000000000005
Run Id: 3
Service Id: 935056
Started By: pateldhwani386@gmail.com
Run Context Type: ICS_UI
06/08/2023 08:28:27 **** Importing Connection: Conn_0123UT0B000000000006 ...
06/08/2023 08:28:27 **** Importing Connection: Conn_0123UT0B000000000007 ...
06/08/2023 08:28:27 **** Importing Source Definition: EMPLOYEES ...
06/08/2023 08:28:27 **** Importing Target Definition: EMPLOYEES ...
06/08/2023 08:28:27 **** Importing SessionConfig: default_session_config ...
  <Warning> : The Error Log DB Connection value should have Relational: as the prefix.
  <Warning> : Invalid value for attribute Error Log DB Connection. Will use the default value
  Validating Source Definition EMPLOYEES...
  Validating Target Definition EMPLOYEES...
06/08/2023 08:28:27 **** Importing Mapping: Mapping0 ...
Validating transformations of mapping Mapping0...
Validating mapping variable(s).
06/08/2023 08:28:27 **** Importing Workflow: wf_mtt_0123UT49000000000005 ...
  <Warning> : Invalid value false for attribute Truncate Target Table. Will use the default value NO
[Session< s_mtt_0123UT49000000000005_0dde8e8d65fb458dae7a7391c64c738e > -->
SnowflakeCloudDataWarehouseV2_table Writer< SnowflakeCloudDataWarehouseV2_table Writer > ]
  <Warning> : Invalid value None for attribute Pushdown Optimization Context. Will use the default value
None [Session< s_mtt_0123UT49000000000005_0dde8e8d65fb458dae7a7391c64c738e > ]
  <Warning> : The value entered is not a valid integer.
  <Warning> : Invalid value NO for attribute Fail task after wait time. Will use the default value
Successfully extracted session instance [s_mtt_0123UT49000000000005_0dde8e8d65fb458dae7a7391c64c738e].
Starting repository sequence id is [1834661625]
DIRECTOR> VAR_27085 [2023-06-08 08:28:27.775] Parameter file
[/data2/home/cldagnt/SystemAgent/apps/Data_Integration_Server/data/0123UT0E00000000005P/0123UT/parameters/s_m
tt_0123UT49000000000005_0dde8e8d65fb458dae7a7391c64c738e.param] is opened for [session
[wf_mtt_0123UT49000000000005.s_mtt_0123UT49000000000005_0dde8e8d65fb458dae7a7391c64c738e]].
DIRECTOR> VAR_27062 [2023-06-08 08:28:27.775] Warning! Cannot find section for worklet
[wf_mtt_0123UT49000000000005] and folder [] in parameter file
[/data2/home/cldagnt/SystemAgent/apps/Data_Integration_Server/data/0123UT0E00000000005P/0123UT/parameters/s_m
tt_0123UT49000000000005_0dde8e8d65fb458dae7a7391c64c738e.param].
DIRECTOR> TM_6014 [2023-06-08 08:28:27.776] Initializing session
[s_mtt_0123UT49000000000005_0dde8e8d65fb458dae7a7391c64c738e] at [Thu Jun 08 08:28:27 2023].
DIRECTOR> TM_6683 [2023-06-08 08:28:27.776] Repository Name: [XMLRepository]

```

```

DIRECTOR> TM_6684 [2023-06-08 08:28:27.776] Server Name: [rDTM]
DIRECTOR> TM_6686 [2023-06-08 08:28:27.776] Folder: []
DIRECTOR> TM_6685 [2023-06-08 08:28:27.776] Workflow: [wf_mtt_0123UT49000000000005] Run Instance Name: [] Run
Id: [0]
DIRECTOR> TM_6992 [2023-06-08 08:28:27.776] Operating System Type [Linux]
DIRECTOR> TM_6101 [2023-06-08 08:28:27.776] Mapping name: Mapping0.
DIRECTOR> TM_6964 [2023-06-08 08:28:27.776] Date format for the Session is [MM/DD/YYYY HH24:MI:SS.US]
DIRECTOR> TM_6703 [2023-06-08 08:28:27.776] Session
[s_mtt_0123UT49000000000005_0dde8e8d65fb458dae7a7391c64c738e] is run by 64-bit Integration Service [],
version [10.2.0 HotFix2], build [0320].
MANAGER> PETL_24091 [2023-06-08 08:28:27.787] Thread [MANAGER] has the ID [1].
MANAGER> PETL_24058 [2023-06-08 08:28:27.787] Running Partition Group [1].
MANAGER> PETL_24000 [2023-06-08 08:28:27.787] Parallel Pipeline Engine initializing.
MONITOR> PETL_24091 [2023-06-08 08:28:27.787] Thread [MONITOR] has the ID [2].
MANAGER> PETL_24001 [2023-06-08 08:28:27.788] Parallel Pipeline Engine running.
MANAGER> PETL_24003 [2023-06-08 08:28:27.788] Initializing session run.
MAPPING> PETL_24091 [2023-06-08 08:28:27.788] Thread [MAPPING] has the ID [3].
MAPPING> CMN_1569 [2023-06-08 08:28:27.790] Server Mode: [UNICODE]
MAPPING> CMN_1570 [2023-06-08 08:28:27.790] Server Code page: [ISO 8859-1 Western European]
MAPPING> TM_6151 [2023-06-08 08:28:27.790] The session sort order is [Binary].
MAPPING> TM_6185 [2023-06-08 08:28:27.790] Warning. Code page validation is disabled in this session.
MAPPING> CMN_65110 [2023-06-08 08:28:27.790] Current Timezone:[UTC -5:0]
MAPPING> CMN_65111 [2023-06-08 08:28:27.790] Current Process ID:[5719]
MAPPING> TM_6155 [2023-06-08 08:28:27.790] Using HIGH precision processing.
MAPPING> TM_6180 [2023-06-08 08:28:27.790] Deadlock retry logic will not be implemented.
MAPPING> TM_6187 [2023-06-08 08:28:27.790] Session target-based commit interval is [10000].
MAPPING> PMJVM_42020 [2023-06-08 08:28:27.793] [INFO] Loaded library :
/data2/home/cldagnt/SystemAgent/jdk/jre/lib/amd64/server/libjvm.so.
MAPPING> PMJVM_42009 [2023-06-08 08:28:27.850] [INFO] Created Java VM successfully.
MAPPING> SDKS_38029 [2023-06-08 08:28:27.900] Loaded plug-in 605501: [plugin description].
MAPPING> SDKS_38024 [2023-06-08 08:28:28.294] Plug-in 605501 initialization complete.
MAPPING> SDKS_38017 [2023-06-08 08:28:28.294] Writer SDK plug-in initialization complete.
MAPPING> CONNECTOR_10000 [2023-06-08 08:28:29.961] [INFO] The source or target object at runtime is
overridden with the parameter name EMPLOYEES= table name EMPLOYEES.
MAPPING> SDKS_38509 [2023-06-08 08:28:30.197] SDK target and group initialization complete.
MAPPING> TM_6307 [2023-06-08 08:28:30.197] DTM error log disabled.
MAPPING> TE_7022 [2023-06-08 08:28:30.197] TShmWriter: Initialized
MAPPING> TM_6007 [2023-06-08 08:28:30.199] DTM initialized successfully for session
[s_mtt_0123UT49000000000005_0dde8e8d65fb458dae7a7391c64c738e]
DIRECTOR> PETL_24033 [2023-06-08 08:28:30.199] All DTM Connection Info: [<NONE>].
MANAGER> PETL_24004 [2023-06-08 08:28:30.200] Starting pre-session tasks. : (Thu Jun 08 08:28:30 2023)
PRE-SESS> PETL_24091 [2023-06-08 08:28:30.200] Thread [PRE-SESS] has the ID [4].
MANAGER> PETL_24027 [2023-06-08 08:28:30.201] Pre-session task completed successfully. : (Thu Jun 08 08:28:30
2023)
DIRECTOR> PETL_24006 [2023-06-08 08:28:30.201] Starting data movement.
MAPPING> TM_6660 [2023-06-08 08:28:30.203] Total Buffer Pool size is 8314384 bytes and Block size is 459688
bytes.
READER_1_1_1> PETL_24091 [2023-06-08 08:28:30.204] Thread [READER_1_1_1] has the ID [5].
READER_1_1_1> DBG_21438 [2023-06-08 08:28:30.204] Reader: Source is
[odbc://dbtype=Oracle;host=testdb.chg7lglng4b.us-east-1.rds.amazonaws.com;port=1521;database=orcl], Type
[ODBC], User [admin]
READER_1_1_1> BLKR_16051 [2023-06-08 08:28:30.207] Source database connection [Conn_0123UT0B000000000006]
code page: [MS Windows Latin 1 (ANSI), superset of Latin1]
READER_1_1_1> CMN_1021 [2023-06-08 08:28:30.207] Database driver event...
CMN_1021 [Driver Manager used is [Data Direct].]

READER_1_1_1> CMN_1021 [2023-06-08 08:28:32.878] Database driver event...
CMN_1021 [Reading driver properties from :
/data2/home/cldagnt/SystemAgent/apps/Data_Integration_Server/65.0.1.1/ICS/main/bin/rdtm/../../../../drivers/d
river.properties
DTM flag to determine driver name is not found hence default to [DataDirect CLOSED 8.0 Oracle Wire Protocol
for Informatica - R40] driver.
ODBC Event Using array fetches.
ODBC Event Using Single Row Inserts. connect string = [Driver=DataDirect CLOSED 8.0 Oracle Wire Protocol for
Informatica - R40;Host=testdb.chg7lglng4b.us-east-
1.rds.amazonaws.com;ServiceName=orcl;Port=1521;UID=admin;AS=6000000;EBL=0;ENS=1;IACP=4;ETWT=1;WorkArounds2=2;
AllowedOpenSSLVersions=1.0.2;EnableEmptyLobs=1]. userid = [admin]
ODBC Version:      03.52.0000
.
DBMS Name:         Oracle
.
DBMS Version:      19.00.0000 Oracle 19.0.0.0.0
.
Driver Name:       DWora28.so
.
Driver Version:    08.02.2432 (B0798, U0577)
.
Driver ODBC Version: 03.52
.]

```

```

READER_1_1_1> BLKR_16003 [2023-06-08 08:28:32.878] Initialization completed successfully.
TRANSF_1_1_1> PETL_24091 [2023-06-08 08:28:32.879] Thread [TRANSF_1_1_1] has the ID [6].
WRITER_1_*_1> PETL_24091 [2023-06-08 08:28:32.879] Thread [WRITER_1_*_1] has the ID [7].
WRITER_1_*_1> CONNECTOR_10000 [2023-06-08 08:28:32.890] [INFO] Normal Mode (Non-Bulk Processing) is enabled
to Write data
WRITER_1_*_1> SNOWFLAKECLOUDDATAWAREHOUSE_10000 [2023-06-08 08:28:36.022] [INFO] The Snowflake Connector uses
the following JDBC URL to connect to Snowflake: jdbc:snowflake://qkwpabt-
gk41830.snowflakecomputing.com/?warehouse=warehouse_target&role=ACCOUNTADMIN&SSL=on&application=INFA_DI_Cloud
WRITER_1_*_1> WRT_8270 [2023-06-08 08:28:37.242] Target connection group #1 consists of target(s) [EMPLOYEES:
Partition 1]
WRITER_1_*_1> WRT_8003 [2023-06-08 08:28:37.242] Writer initialization complete.
WRITER_1_*_1> WRT_8005 [2023-06-08 08:28:37.242] Writer run started.
READER_1_1_1> BLKR_16007 [2023-06-08 08:28:37.242] Reader run started.
WRITER_1_*_1> WRT_8158 [2023-06-08 08:28:37.242]

```

*****START LOAD SESSION*****

Load Start Time: Thu Jun 08 08:28:36 2023

Target tables:

EMPLOYEES: Partition 1

```

READER_1_1_1> RR_4010 [2023-06-08 08:28:37.243] SQ instance [Source] SQL Query [SELECT
ADMIN.EMPLOYEES.EMPLOYEE_ID, ADMIN.EMPLOYEES.FIRST_NAME, ADMIN.EMPLOYEES.LAST_NAME, ADMIN.EMPLOYEES.EMAIL,
ADMIN.EMPLOYEES.PHONE, ADMIN.EMPLOYEES.HIRE_DATE, ADMIN.EMPLOYEES.MANAGER_ID, ADMIN.EMPLOYEES.JOB_TITLE FROM
ADMIN.EMPLOYEES]
READER_1_1_1> RR_4049 [2023-06-08 08:28:37.243] SQL Query issued to database : (Thu Jun 08 08:28:37 2023)
READER_1_1_1> RR_4050 [2023-06-08 08:28:37.668] First row returned from database to reader : (Thu Jun 08
08:28:37 2023)
READER_1_1_1> BLKR_16019 [2023-06-08 08:28:37.669] Read [107] rows, read [0] error rows for source table
[EMPLOYEES] instance name [EMPLOYEES]
READER_1_1_1> BLKR_16008 [2023-06-08 08:28:37.669] Reader run completed.
TRANSF_1_1_1> DBG_21216 [2023-06-08 08:28:37.669] Finished transformations for Source Qualifier [Source].
Total errors [0]
WRITER_1_*_1> WRT_8167 [2023-06-08 08:28:37.669] Start loading table [EMPLOYEES: Partition 1] at: Thu Jun 08
08:28:36 2023
WRITER_1_*_1> WRT_8168 [2023-06-08 08:28:37.693] End loading table [EMPLOYEES: Partition 1] at: Thu Jun 08
08:28:36 2023
WRITER_1_*_1> SNOWFLAKECLOUDDATAWAREHOUSE_10000 [2023-06-08 08:28:37.693] [INFO] The Snowflake Connector
completed writing data to the target.
WRITER_1_*_1> WRT_8035 [2023-06-08 08:28:45.200] Load complete time: Thu Jun 08 08:28:44 2023

```

LOAD SUMMARY

=====

```

WRT_8036 Target: EMPLOYEES: Partition 1 (Instance Name: [EMPLOYEES])
WRT_8038 Inserted rows - Requested: 107          Applied: 107          Rejected: 0          Affected: 107

```

```

WRITER_1_*_1> WRT_8043 [2023-06-08 08:28:45.200] *****END LOAD SESSION*****
WRITER_1_*_1> WRT_8006 [2023-06-08 08:28:45.200] Writer run completed.
MANAGER> PETL_24031 [2023-06-08 08:28:45.201]
***** RUN INFO FOR TGT LOAD ORDER GROUP [1], CONCURRENT SET [1] *****
Thread [READER_1_1_1] created for [the read stage] of partition point [Source] has completed.
    Total Run Time = [0.627216] secs
    Total Idle Time = [0.000000] secs
    Busy Percentage = [100.000000]
Thread [TRANSF_1_1_1] created for [the transformation stage] of partition point [Source] has completed. The
total run time was insufficient for any meaningful statistics.
Thread [WRITER_1_*_1] created for [the write stage] of partition point [EMPLOYEES] has completed.
    Total Run Time = [7.531312] secs
    Total Idle Time = [0.000000] secs
    Busy Percentage = [100.000000]

MANAGER> PETL_24005 [2023-06-08 08:28:45.202] Starting post-session tasks. : (Thu Jun 08 08:28:45 2023)
POST-SESS> PETL_24091 [2023-06-08 08:28:45.202] Thread [POST-SESS] has the ID [8].
MANAGER> PETL_24029 [2023-06-08 08:28:45.202] Post-session task completed successfully. : (Thu Jun 08
08:28:45 2023)
MAPPING> SDKS_38510 [2023-06-08 08:28:45.203] SDK target and group deinitialized with status [0].
MAPPING> SDKS_38025 [2023-06-08 08:28:45.203] Plug-in 605501 deinitialized and unloaded with status [0].
MAPPING> SDKS_38019 [2023-06-08 08:28:45.203] Writer SDK plug-ins deinitialized with status [0].
MAPPING> TM_6018 [2023-06-08 08:28:45.203] The session completed with [0] row transformation errors.
MANAGER> PETL_24002 [2023-06-08 08:28:45.897] Parallel Pipeline Engine finished.
DIRECTOR> PETL_24012 [2023-06-08 08:28:45.897] Session run completed successfully.
DIRECTOR> TM_6022 [2023-06-08 08:28:45.897]

```


SESSION LOAD SUMMARY

=====

DIRECTOR> TM_6252 [2023-06-08 08:28:45.897] Source Load Summary.

DIRECTOR> CMN_1740 [2023-06-08 08:28:45.897] Table: [Source] (Instance Name: [Source] Instance UI Name: [Source])

Output Rows [107], Affected Rows [107], Applied Rows [107], Rejected Rows [0]

DIRECTOR> TM_6253 [2023-06-08 08:28:45.897] Target Load Summary.

DIRECTOR> CMN_1740 [2023-06-08 08:28:45.897] Table: [EMPLOYEES] (Instance Name: [EMPLOYEES] Instance UI Name: [Target])

Output Rows [107], Affected Rows [107], Applied Rows [107], Rejected Rows [0]

DIRECTOR> TM_6023 [2023-06-08 08:28:45.897]

=====

DIRECTOR> TM_6020 [2023-06-08 08:28:45.897] Session

[s_mtt_0123UT49000000000005_0dde8e8d65fb458dae7a7391c64c738e] completed at [Thu Jun 08 08:28:45 2023].